

DYNAMIC CONTEXT-AWARE TRIES FOR KNOWLEDGE GRAPH-CONSTRAINED LARGE LANGUAGE MODEL GENERATION

BERNARD KIRK ADJANOR KATAMANSO
ERICA AMONOR
JOSEPH OSEI NYARKO
JESSICA AFUA ETORNAM NSAFOAH

Supervised By: DR ERIC AFFUM

University of Mines and Technology
Faculty of Computing and Mathematical Science
Computer Science and Engineering Department

July ,2026



Outline

- 1 Introduction
- 2 Problem Statement
- 3 Objectives
- 4 Methodology
- 5 Results
- 6 Discussion & Contributions
- 7 Recommendations
- 8 References



Introduction

- LLMs hallucinate, producing factually incorrect outputs
- KGQA requires answers that follow real graph structure
- Constrained decoding (GCR, DoG) guarantees validity but not relevance
- A 3-hop question on Freebase can produce $\sim 24,000$ structurally valid paths, yet only one is correct



Statement of Problem

- GCR's trie is fixed before decoding starts, so it cannot adapt to question intent or the evolving chain
- The LLM must filter semantically irrelevant paths using internal knowledge, the same mechanism that causes hallucinations



Project Objectives (1/2)

- Characterize the permissiveness problem by measuring SIR of GCR's static KG-Tries on WebQSP
- Design a symbolic constraint oracle (TypeOracle) with range and type gates, avoiding embedding cost and threshold sensitivity



Project Objectives (2/2)

- Implement DCA-Trie v1 (static filtering at construction time, $\text{FNR} < 5\%$) and v2 (dynamic expansion per committed triplet)
- Evaluate both against GCR baseline on WebQSP: Hits@1, F1, structural faithfulness, SIR, trie size



Software

- Python, PyTorch, HuggingFace
- GCR Framework
- Sentence-transformers (MiniLM)
- NetworkX, marisa_trie
- Freebase (WebQSP, CWQ)

Hardware & Sources

- NVIDIA RTX 4090 (24 GB)
- ACM, arXiv, ACL Anthology



- Literature review and theoretical baseline
- Algorithm design and TypeOracle implementation
- Experimental benchmarking on WebQSP and CWQ



Methodology




Figure: DCA-Trie System Architecture.



Baseline (GCR/DoG):

$$W_{\text{val}}^{\text{GCR}}(t) = f(G, E_q)$$

Proposed (DCA-Trie):

$$W_{\text{val}}^{\text{DCA}}(t) = f(G, E_q, q, y_{<t})$$



- **F (Faithfulness):** 100% structural grounding in KG triplets
- **R (Relevance):** Lower SIR than baseline
- **P (Recall):** FNR on gold paths $< 5\%$



Constrained Decoding in LLMs

At each token step t , the LLM produces logits $\mathbf{h}_t \in \mathbb{R}^{|V|}$ over the vocabulary V . A constraint mask $\mathbf{m}_t \in \{0, 1\}^{|V|}$ zeros out invalid tokens:

$$\tilde{P}(x_t) = \text{softmax}(\mathbf{h}_t + (1 - \mathbf{m}_t) \cdot (-\infty))$$

- In GCR/DoG, $\mathbf{m}_t$ comes from a static trie built before decoding:
 $\mathbf{m}_t = f(G, E_q)$
- DCA-Trie conditions the mask on the question and partial output:
 $\mathbf{m}_t = f(G, E_q, q, y_{<t})$
- Only tokens corresponding to valid KG triples receive non-zero probability



TypeOracle: Gate Design

Two symbolic gates over KG ontology metadata, applied per candidate triple (e_t, r, e') :

Gate 1: Property Range Gate (every hop)

$$\text{range_gate}(r, e') = \mathbf{1}[\text{types}(e') \cap \text{range}(r) \neq \emptyset]$$

Checks that the tail entity's types fall within the relation's expected range.

Gate 2: Answer Type Gate (terminal hop only)

$$\text{type_gate}(e', q, h, L) = \mathbf{1}[\text{types}(e') \cap \mathcal{T}(q) \neq \emptyset]$$

Checks that the terminal entity matches the inferred answer type $\mathcal{T}(q)$.

Combined: $\text{admissible} = \text{range_gate} \wedge \text{type_gate}$, with conservative fallback (admit when schema info is missing).



Diagnostic Metric: SIR

SIR measures the fraction of structurally valid paths that are semantically incompatible with the question.

$$\text{SIR}(q, t) = \frac{\sum_{p \in \mathcal{P}(q, t)} \text{irrelevant}(p, q)}{|\mathcal{P}(q, t)|}$$

Corpus average:

$$\overline{\text{SIR}} = \frac{1}{|Q|} \sum_{q \in Q} \frac{1}{L_q} \sum_{t=1}^{L_q} \text{SIR}(q, t)$$



- Filters BFS-enumerated paths through TypeOracle before building the prefix tree
- Preserves $O(d)$ trie lookup while pruning irrelevant branches offline




DCA-Trie v2: Dynamic Expansion

- Expands the trie step-by-step as entity tokens e_t are committed
- Applies range and type gate checks to $\mathcal{N}(e_t)$ during generation




Results: Hits@1 Comparison



WebQSP, 1,628 test queries, Llama-3.1-8B, beam $k=10$



Results: Constraint Quality



- 14.5% path reduction (4.10M $\rightarrow$ 3.51M paths)
- FNR: 3.3% (type gate), 2.9% (range gate)
- Latency overhead: +0.5%



Key Contributions

Contribution	Before	After	Why It Matters
TypeOracle	Embedding + threshold τ	Two symbolic gates	No encoder, no τ , $O(1)$, deterministic
SIR Metric	Accuracy only	Semantic Irrelevance Ratio	Diagnoses constraint quality independently
Non-Monotone Finding	"Tighter $\rightarrow$ better" (GCR, DoG assumption)	Tightness $\neq$ accuracy ($\Delta\text{Acc} = -5.2\text{pp}$)	First evidence constraint tightness is not a proxy for quality



- **Non-monotone:** Tighter constraints do not guarantee better answers.
- **Four degradation mechanisms:** gold-path exclusion, beam competition, regex type-inference noise, precision–recall trade-off.
- **TypeOracle efficiency:** $O(1)$ per path, no encoder, no threshold, only +0.5% latency.
- **SIR fills a gap:** diagnoses constraint quality independently of answer accuracy.



Live demo: GCR baseline vs DCA-Trie v1
on a sample WebQSP question



Conclusion

- DCA-Trie v1 meets all three conditions: 100% faithfulness, 14.5% path reduction, $\text{FNR} < 5\%$.
- TypeOracle achieves meaningful tightening with $O(1)$ symbolic gates.
- Central finding: constraint tightness and accuracy are **non-monotone**, with Hits@1 dropping 5.2pp.
- SIR is introduced as the first diagnostic for constraint quality, decoupled from accuracy.



Recommendations

- Replace regex type classifier with LLM-based intent parser (85% $\rightarrow$ higher recall).
- Extend TypeOracle schema beyond 40 Freebase relations.
- Explore hybrid oracles (symbolic + embedding) to recover beam diversity.
- Evaluate on CWQ with $N \geq 500$ to reduce variance.



Thank You!

Thank You!!!



References (1/2)

- Brown, T. et al. (2020). Language Models are Few-Shot Learners. *Advances in NeuralPS*, 33, 1877–1901.
- De Cao, N. et al. (2021). Autoregressive Entity Retrieval. *ICLR*.
- Geng, S. et al. (2023). Grammar-Constrained Decoding for Structured NLP Tasks without Finetuning. *Proceedings of EMNLP*.
- He, G. et al. (2021). Improving Multi-Hop KBQA by Learning Intermediate Supervision Signals. *Proceedings of WSDM*, 553–556.



References (2/2)

- Lan, Y. et al. (2023). Complex Knowledge Base Question Answering: A Survey. *IEEE Trans. Knowledge and Data Engineering*, 35(11), 11196–11215.
- Lewis, P. et al. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. *Advances in NeurIPS*, 33, 9459–9474.
- Li, K. et al. (2025). Decoding on Graphs: Faithful and Sound Reasoning on Knowledge Graphs. *Proceedings of ACL*, 24349–24364.
- Luo, L. et al. (2025). Graph-Constrained Reasoning: Faithful Reasoning on KGs with LLMs. *Proceedings of ICML*, 267, 41540–41565.
- Willard, B. T. and Louf, R. (2023). Efficient Guided Generation for Large Language Models. *arXiv:2307.09702*.

